

# **COMMONWEALTH OF AUSTRALIA**

## **Copyright Regulations 1969**

### **WARNING**

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**). The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

**Do not remove this notice.**

# COMP2823

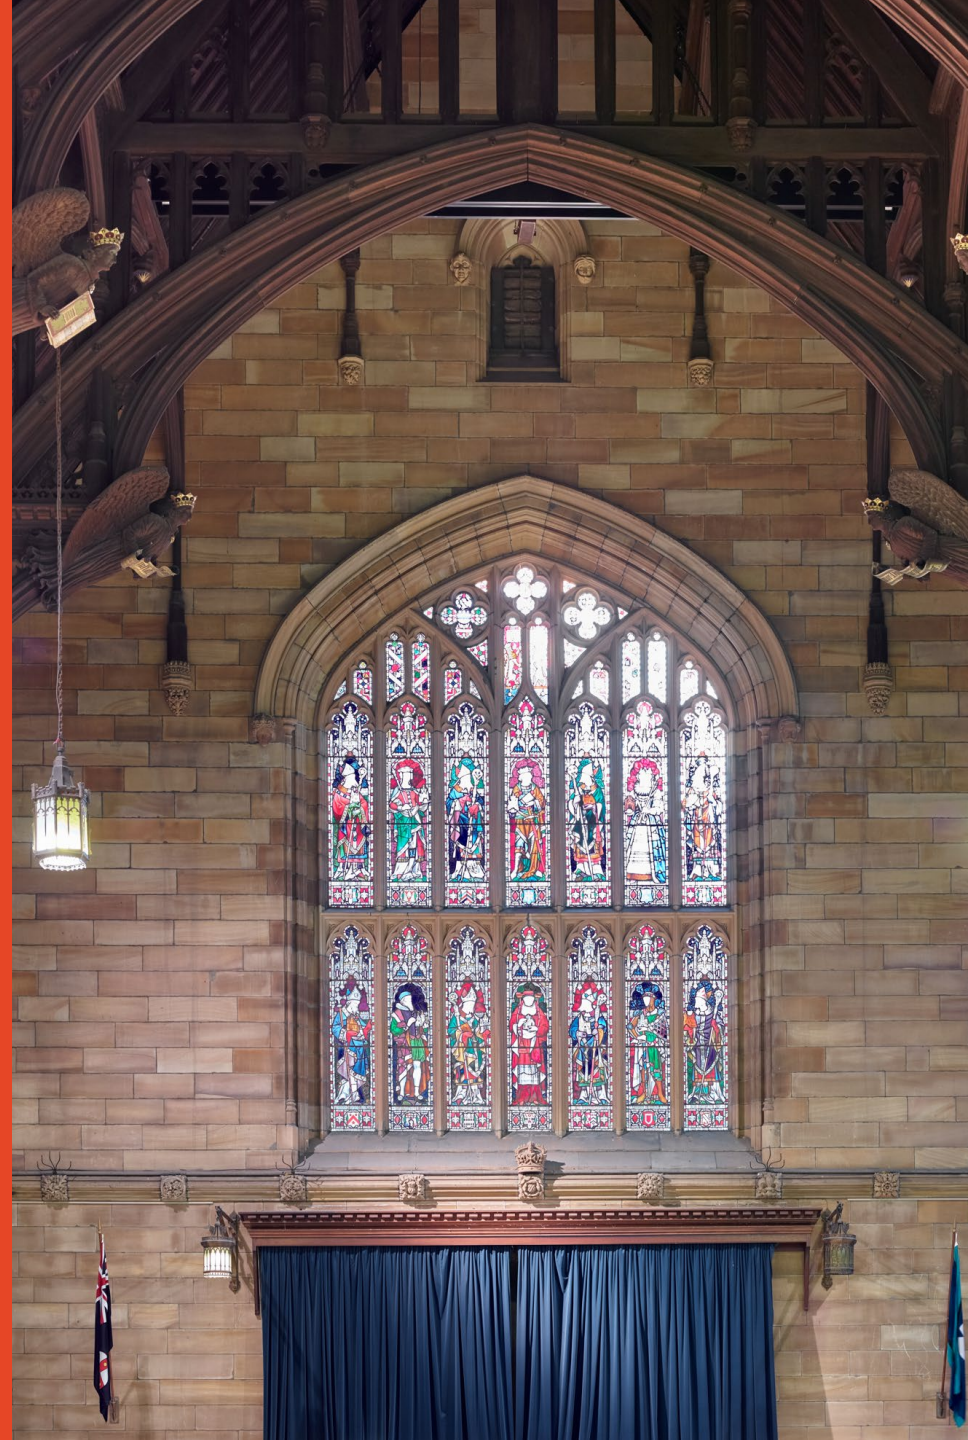
## Lecture 5: Priority Queues [GT 5]

Joachim Gudmundsson  
School of Computer Science

*Some content is taken from material  
provided by the textbook publisher Wiley.*



THE UNIVERSITY OF  
SYDNEY



# Priority Queue ADT

Special type of ADT map to store a collection of key-value items where we can only remove smallest key:

- **insert**(*k*, *v*): insert item with key *k* and value *v*
- **remove\_min**(): remove and return the item with smallest key
- **min**(): return item with smallest key
- **size**(): return how many items are stored
- **is\_empty**(): test if queue is empty

We can also have a max version of this min version, but we need one structure for each version.

# Example

A sequence of priority queue methods:

Method	Return value	Priority queue
insert(5,A)		{(5,A)}
insert(9,C)		{(5,A),(9,C)}
insert(3,B)		{(3,B),(5,A),(9,C)}
min()	(3,B)	{(3,B),(5,A),(9,C)}
remove_min()	(3,B)	{(5,A),(9,C)}
insert(7,D)		{(5,A),(7,D),(9,C)}
remove_min()	(5,A)	{(7,D),(9,C)}
remove_min()	(7,D)	{(9,C)}
remove_min()	(9,C)	{}
is_empty()	true	{}

# Application: Stock Matching Engines

At the heart of modern stock trading systems are highly reliable systems known as **matching engines**, which match the stock trades of buyers and sellers.

Buyers post bids to buy a number of shares of a given stock at or below a specified price


Sellers post offers (asks) to sell a number of shares of a given stock at or above a specified price.

STOCK: EXAMPLE.COM					
Buy Orders			Sell Orders		
Shares	Price	Time	Shares	Price	Time
1000	4.05	20 s	500	4.06	13 s
100	4.05	6 s	2000	4.07	46 s
2100	4.03	20 s	400	4.07	22 s
1000	4.02	3 s	3000	4.10	54 s
2500	4.01	81 s	500	4.12	2 s
			3000	4.20	58 s
			800	4.25	33 s
			100	4.50	92 s

# Application: Stock Matching Engines

Buy and sell orders are organized according to a **price-time priority**, where price has highest priority and time is used to break ties

When a new order is entered, the matching engine determines if a trade can be immediately executed and if so, then it performs the appropriate matches according to price-time priority.



## STOCK: EXAMPLE.COM

### Buy Orders

### Sell Orders

Shares	Price	Time	Shares	Price	Time
1000	4.05	20 s	500	4.06	13 s
100	4.05	6 s	2000	4.07	46 s
2100	4.03	20 s	400	4.07	22 s
1000	4.02	3 s	3000	4.10	54 s
2500	4.01	81 s	500	4.12	2 s
			3000	4.20	58 s
			800	4.25	33 s
			100	4.50	92 s

# Application: Stock Matching Engines

A matching engine can be implemented with two **priority queues**, one for buy orders and one for sell orders.

This data structure performs element removals based on priorities assigned to elements when they are inserted.

```
while True:
    bid ← buy_orders.max()
    ask ← sell_orders.min()
    if bid.price ≥ ask.price then
        bid ← buy_orders.remove_max()
        ask ← sell_orders.remove_min()
        carry out trade (bid, ask)
```

STOCK: EXAMPLE.COM					
Buy Orders			Sell Orders		
Shares	Price	Time	Shares	Price	Time
1000	4.05	20 s	500	4.06	13 s
100	4.05	6 s	2000	4.07	46 s
2100	4.03	20 s	400	4.07	22 s
1000	4.02	3 s	3000	4.10	54 s
2500	4.01	81 s	500	4.12	2 s
			3000	4.20	58 s
			800	4.25	33 s
			100	4.50	92 s

**Other applications:** sorting, shortest path, task scheduling, data compression, routing, and many more.

# Sequence-based Priority Queue

## Unsorted list implementation



- **insert** in  $O(1)$  time since we can insert the item at the beginning or end of the sequence
- **remove\_min** and **min** in  $O(n)$  time since we have to traverse the entire list to find the smallest key

## Sorted list implementation



- **insert** in  $O(n)$  time since we have to find the place where to insert the item
- **remove\_min** and **min** in  $O(1)$  time since the smallest key is at the beginning

Method	Unsorted List	Sorted List
size, isEmpty	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$
min, removeMin	$O(n)$	$O(1)$




# BST-based Priority Queue

– insert in  $O(\log n)$  time

– remove\_min and min in  $O(\log n)$  time


Method	Unsorted List	Sorted List	BST
size, isEmpty	$O(1)$	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$	$O(\log n)$
min, removeMin	$O(n)$	$O(1)$	$O(\log n)$



# Heap data structure (min-heap)

A **heap** is a binary tree storing (key, value) items at its nodes, satisfying the following properties:

1. **Heap-Order:** for every node  $v \neq \text{root}$ ,  
 $\text{key}(m) \geq \text{key}(\text{parent}(v))$




2. **Complete Binary Tree:** let  $h$  be the height
  - every level  $i < h$  is full (i.e., there are  $2^i$  nodes)
  - remaining nodes take leftmost positions of level  $h$

The **last node** is the rightmost node of maximum depth

# Example

1. **Heap-Order:** for every node  $m \neq \text{root}$ ,  $\text{key}(m) \geq \text{key}(\text{parent}(m))$
2. **Complete Binary Tree:** left to right




# Minimum of a Heap

**Fact:** The root always holds the smallest key in the heap

**Proof:**

- Suppose the minimum key is at some internal node  $x$
- Because of the heap property, as we move up the tree, the keys can only get smaller (assuming repeats, otherwise contradiction)
- If  $x$  is not the root, then its parent must also hold a smallest key
- Keep going until we reach the root




# Height of a Heap

**Fact:** A heap storing  $n$  keys has height  $\log n$

**Proof:**

- Let  $h$  be the height of a heap storing  $n$  keys
- Since there are  $2^i$  keys at depth  $i = 0, \dots, h - 1$  and at least one key at depth  $h$ , we have  $n \geq 1 + 2 + 4 + \dots + 2^{h-1} + 1 = 2^h$
- Thus,  $n \geq 2^h$ , applying  $\log_2$  on both sides,  $\log_2 n \geq h$




# Insertion into a Heap


- a) Create a new node with given key
- b) Find location for new node (new last node)
- c) Restore the heap-order property

insert(1)

a)



b)




# Upheap

Restore heap-order property by swapping keys along upward path from insertion point

```
def up_heap(z):  
    while z ≠ root and key(parent(z)) > key(z) do  
        swap key of z and parent(z)  
        z ← parent(z)
```

Correctness: after swapping the subtree rooted at **z** has the heap-order property

Complexity:  $O(\log n)$  time because the height of the heap is  $\log n$




# Finding the position for insertion

Two possible solutions:





1. Keep track of size of subtrees, or
2. Search starting from last node
  - a) go up until a left child or the root is reached
  - b) If we reach the root then need to open a new level
  - c) otherwise, go to the sibling (right child of parent)
  - d) go down left until a leaf is reached

Complexity of this search is  $O(\log n)$  because the height is  $\log n$ . Thus, overall complexity of insertion is  $O(\log n)$  time











## Example insertion (2,T)



## Example insertion (2,T) cont'd





# Example insertion



# Removal from a Heap

- a) Replace the root key with the key of the last node  $w$
- b) Delete  $w$
- c) Restore the heap-order property

remove\_min()






# Downheap

Restore heap-order property by swapping keys along downward path from the root





```
def down_heap(z):  
    while z has child with  
        key(child) < key(z) do  
        x ← child of z with smallest key  
        swap keys of x and z  
        z ← x
```

Correctness: after swap **z** heap-order property is restored up to level of **z**





Complexity:  $O(\log n)$  time because the height of the heap is  $\log n$





# Example removal



## Example removal



## Example removal






# Finding next last node after deletion

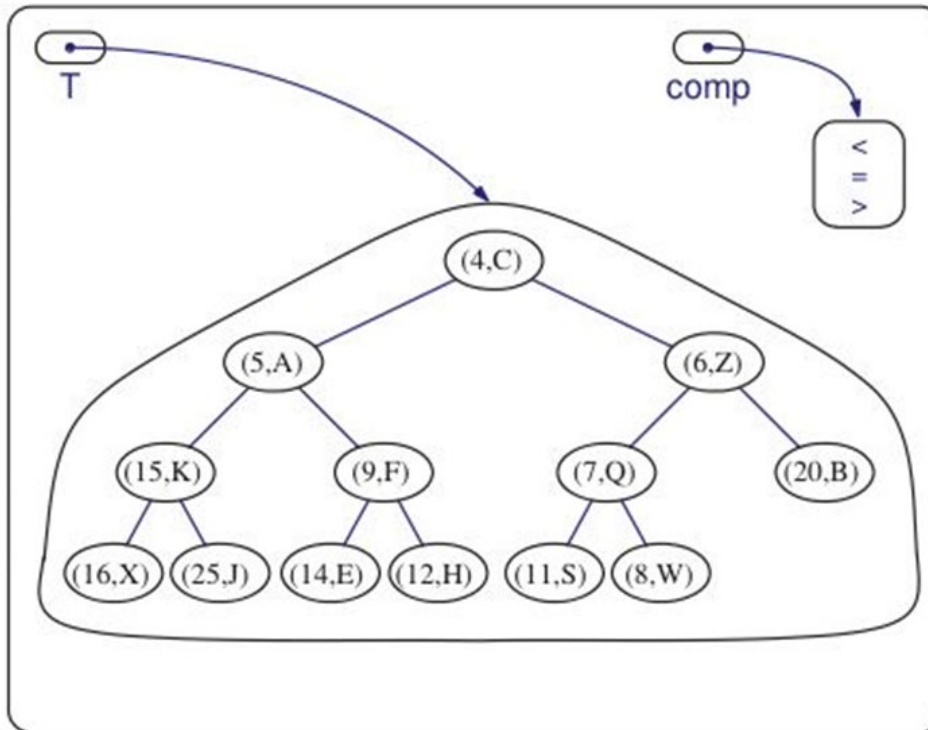
Two possible solutions:

1. Keep track of size of subtrees, or
2. start from the (old) last node
  - a) go up until a right child or the root is reached
  - b) if we reach the root then need to close a level
  - c) otherwise, go to the sibling (left child of parent)
  - d) go down right until a leaf is reached

Complexity of this search is  $O(\log n)$  because the height is  $\log n$ . Thus, overall complexity of deletion is  $O(\log n)$  time





# Heap-based implementation of a priority queue



Operation	Time
size, isEmpty	$O(1)$
min,	$O(1)$
insert	$O(\log n)$
removeMin	$O(\log n)$

# Heap implementation

Heaps can be implemented as a linked structure or an array.



2	5	6	9	7
0	1	2	3	4

# Heap-in-array implementation



We can represent a heap with  $n$  keys by means of an array of length  $n$

Special nodes:

- root is at  $0$
- last node is at  $n-1$

For the node at index  $i$ :

- the left child is at index  $2i+1$
- the right child is at index  $2i+2$
- Parent is at index  $\lfloor (i-1)/2 \rfloor$



## Summary: Priority queue implementations

Method	Unsorted List	Sorted List	Heap	BST
size, isEmpty	$O(1)$	$O(1)$	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$	$O(\log n)$	$O(\log n)$
min	$O(n)$	$O(1)$	$O(1)$	$O(\log n)$
removeMin	$O(n)$	$O(1)$	$O(\log n)$	$O(\log n)$

# Building a heap in one go

Sometimes we have all the keys upfront. If we insert them one at a time, this can take  $O(n \log n)$  time. However, there is a faster way to build the heap in this case.

```
def heapify (A):  
    # turn A into a binary heap in place  
    n ← size(A)  
    for i from n-1 to 0 do  
        down_heap(A, i)
```

If we let  $h(i)$  be the height of the node corresponding to  $A[i]$  then  $\text{down\_heap}(A, i)$  takes  $O(h(i))$  time.

Thus, the running time of the algorithm is  $O(\sum_{i=0}^{n-1} h(i))$




# Downheap

Restore heap-order property by swapping keys along downward path from the root


```
def down_heap(z):  
    while z has child with  
        key(child) < key(z) do  
        x ← child of z with smallest key  
        swap keys of x and z  
        z ← x
```

Correctness: after swap **z** heap-order property is restored up to level of **z**

Complexity:  $O(\log n)$  time because the height of the heap is  $\log n$




# Building a heap in one go







# Building a heap in one go




# Building a heap in one go




# Building a heap in one go




# Building a heap in one go




# Building a heap in one go




# Building a heap in one go



# Building a heap in one go




# Building a heap in one go






# Building a heap in one go



# Building a heap in one go




# Building a heap in one go

**Claim:** The running time of the algorithm is  $O(\sum_{i=0}^{n-1} h(i)) = O(n)$

For each node  $i$  in the tree construct a path of length  $h(i)$  by starting at node  $i$ , going **right once**, and then going left until we reach a leaf.

**Claim:** These paths are edge disjoint




# Building a heap in one go

**Claim:** The running time of the algorithm is  $O(\sum_{i=0}^{n-1} h(i)) = O(n)$

For each node  $i$  in the tree construct a path of length  $h(i)$  by starting at node  $i$ , going **right once**, and then going left until we reach a leaf.

**Claim:** These paths are edge disjoint




# Building a heap in one go

**Claim:** The running time of the algorithm is  $O(\sum_{i=0}^{n-1} h(i)) = O(n)$

For each node  $i$  in the tree construct a path of length  $h(i)$  by starting at node  $i$ , going **right once**, and then going left until we reach a leaf.

**Claim:** These paths are edge disjoint




## Building a heap in one go

**Claim:** The running time of the algorithm is  $O(\sum_{i=0}^{n-1} h(i)) = O(n)$

For each node  $i$  in the tree construct a path of length  $h(i)$  by starting at node  $i$ , going **right once**, and then going left until we reach a leaf.

**Claim:** These paths are edge disjoint




# Building a heap in one go

**Claim:** The running time of the algorithm is  $O\left(\sum_{i=0}^{n-1} h(i)\right) = O(n)$

For each node  $i$  in the tree construct a path of length  $h(i)$  by starting at node  $i$ , going **right once**, and then going left until we reach a leaf.

**Claim:** These paths are edge disjoint




## Building a heap in one go

**Claim:** The running time of the algorithm is  $O(\sum_{i=0}^{n-1} h(i)) = O(n)$

For each node  $i$  in the tree construct a path of length  $h(i)$  by starting at node  $i$ , going **right once**, and then going left until we reach a leaf.

**Claim:** These paths are edge disjoint






# Building a heap in one go

**Claim:** The running time of the algorithm is  $O(\sum_{i=0}^{n-1} h(i)) = O(n)$

For each node  $i$  in the tree construct a path of length  $h(i)$  by starting at node  $i$ , going **right once**, and then going left until we reach a leaf.

**Claim:** These paths are edge disjoint




## Building a heap in one go

**Claim:** The running time of the algorithm is  $O(\sum_{i=0}^{n-1} h(i)) = O(n)$


For each node  $i$  in the tree construct a path of length  $h(i)$  by starting at node  $i$ , going **right once**, and then going left until we reach a leaf.

**Claim:** These paths are edge disjoint



# Sequence-based Priority Queue

## Unsorted list implementation



- **insert** in  $O(1)$  time since we can insert the item at the beginning or end of the sequence
- **remove\_min** and **min** in  $O(n)$  time since we have to traverse the entire list to find the smallest key

## Sorted list implementation



- **insert** in  $O(n)$  time since we have to find the place where to insert the item
- **remove\_min** and **min** in  $O(1)$  time since the smallest key is at the beginning


Method	Unsorted List	Sorted List
size, isEmpty	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$
min, removeMin	$O(n)$	$O(1)$

# BST-based Priority Queue

– insert in  $O(\log n)$  time

– remove\_min and min in  $O(\log n)$  time

Method	Unsorted List	Sorted List	BST
size, isEmpty	$O(1)$	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$	$O(\log n)$
min, removeMin	$O(n)$	$O(1)$	$O(\log n)$




# Heap-based Priority Queue

insert in  $O(\log n)$  time

remove\_min and min in  $O(\log n)$  time

Can be built in  $O(n)$  time

Summary:



Method	Unsorted List	Sorted List	Heap	BST
size, isEmpty	$O(1)$	$O(1)$	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$	$O(\log n)$	$O(\log n)$
min	$O(n)$	$O(1)$	$O(1)$	$O(\log n)$
removeMin	$O(n)$	$O(1)$	$O(\log n)$	$O(\log n)$

# Priority Queue Sorting

We can use a priority queue to sort a list of keys:

1. iteratively insert keys into an empty priority queue
2. iteratively **remove\_min** to get the keys in sorted order

Complexity analysis:

- **n** insert operations
- **n** remove\_min operations

Sorting using **a**:

List  $O(n^2)$

BST  $O(n \log n)$

Heaps  $O(n \log n)$

```
def priority_queue_sorting(A):  
    pq ← new priority queue  
    n ← size(A)  
    for i in [0:n] do  
        pq.insert(A[i])  
    for i in [0:n] do  
        A[i] ← pq.remove_min()
```

Method	Unsorted List	Sorted List	Heap	BST
size, isEmpty	$O(1)$	$O(1)$	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$	$O(\log n)$	$O(\log n)$
min	$O(n)$	$O(1)$	$O(1)$	$O(\log n)$
removeMin	$O(n)$	$O(1)$	$O(\log n)$	$O(\log n)$

# Selection-Sort

Variant of pq-sort using unsorted sequence implementation:

1. inserting elements with  $n$  insert operations takes  $O(n)$  time
2. removing elements with  $n$  remove\_min operations takes  $O(n^2)$

Can be done in place  
(no need for extra space)

Top level loop invariant:

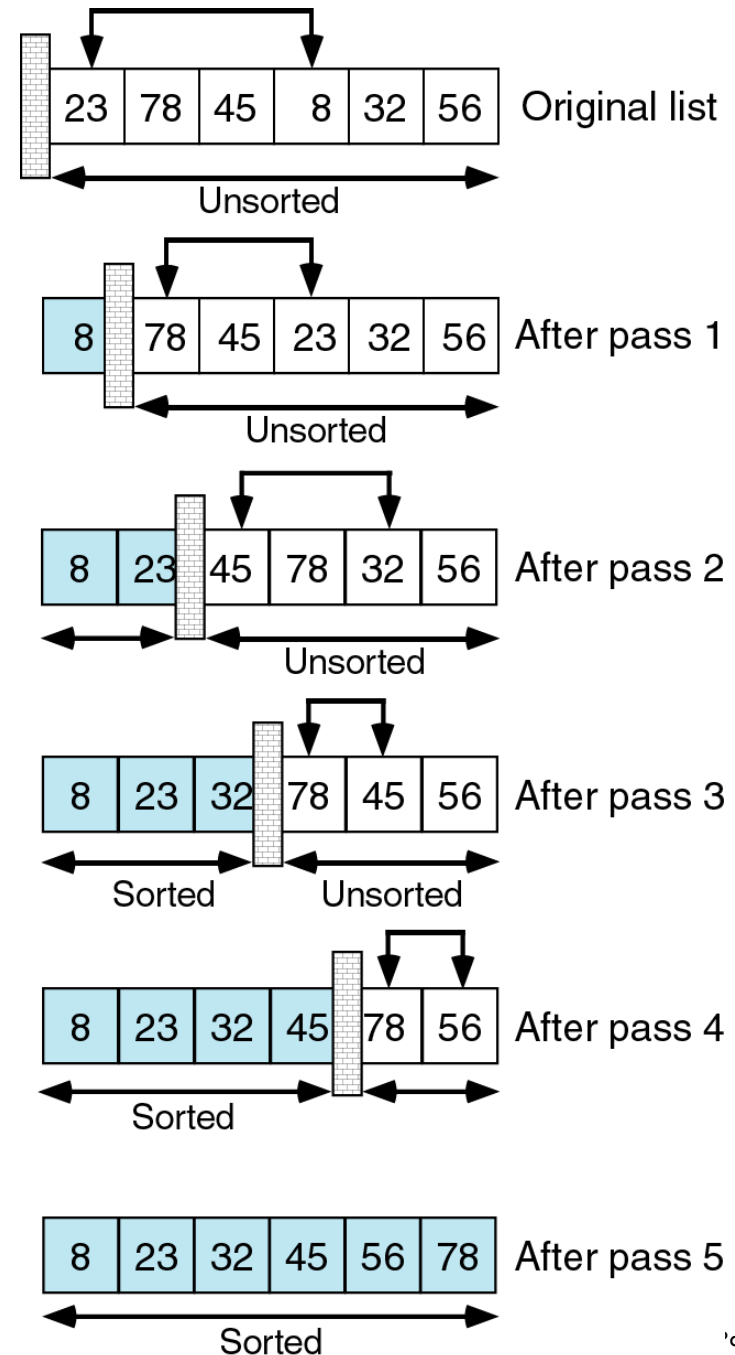
- $A[0:i]$  is sorted
- $A[i:n]$  is the priority queue  
and all  $\geq A[i-1]$

```
def selection_sort(A):  
    n ← size(A)  
    for i in [0:n] do  
        # find s ≥ i minimizing A[s]  
        s ← i  
        for j in [i+1:n] do  
            if A[j] < A[s] then  
                s ← j  
        # swap A[i] and A[s]  
        A[i], A[s] ← A[s], A[i]
```

# Selection-Sort Example

```
def selection_sort(A):  
    n ← size(A)  
    for i in [0:n] do  
        s ← i  
        for j in [i+1:n] do  
            if A[j] < A[s] then  
                s ← j  
        A[i], A[s] ← A[s], A[i]
```

Running time:  $O(n^2)$





# Insertion-Sort

Variant of pq-sort using sorted sequence implementation:

1. inserting elements with  $n$  insert operations takes  $O(n^2)$  time
2. removing elements with  $n$  remove\_min operations takes  $O(n)$

Can be done in place  
(no need for extra space)

Top level loop invariant:

- $A[0:i]$  is the priority queue  
(and thus sorted)
- $A[i:n]$  is yet-to-be-inserted


```
def insertion_sort(A):  
    n ← size(A)  
    for i in [1:n] do  
        x ← A[i]  
        # move forward entries > x  
        j ← i  
        while j > 0 and x < A[j-1] do  
            A[j] ← A[j-1]  
            j ← j - 1  
        # if j>0 ⇒ x ≥ A[j-1]  
        # if j<i ⇒ x < A[j+1]  
        A[j] ← x
```

# Insertion-Sort Example

```
def insertion_sort(A):  
    n ← size(A)  
    for i in [1:n] do  
        x ← A[i]  
        j ← i  
        while j > 0 and x < A[j-1] do  
            A[j] ← A[j-1]  
            j ← j - 1  
        A[j] ← x
```

**Running time:**  $O(n^2)$

- What if the input is sorted?
- What if the input is sorted in descending order?



# Heap-Sort

Consider a priority queue with  $n$  items implemented with a heap:

- the space used is  $O(n)$
- methods **insert** and **remove\_min** take  $O(\log n)$

Recall that priority-queue sorting uses:

- $n$  insert ops
- $n$  remove\_min ops

The  $n$  insertions can be done in one go in  $O(n)$  time.

**Heap-sort** is the version of priority-queue sorting that implements the priority queue with a heap. It runs in  $O(n \log n)$  time.

# Refinements and Generalization

Heap-sort can be arranged to work in place using part of the array for the output and part for the priority queue

A heap on  $n$  keys can be constructed in  $O(n)$  time. But the  $n$  `remove_min` still take  $O(n \log n)$  time

Sometimes it is useful to support a few more operations (all given a pointer to  $e$ ):

- `remove(e)`: Remove item  $e$  from the priority queue
- `replace_key(e, k)`: update key of item  $e$  with  $k$
- `replace_value(e, v)`: update value of item  $e$  with  $v$

Method	Unsorted List	Sorted List	Heap	BST
size, isEmpty	$O(1)$	$O(1)$	$O(1)$	$O(1)$
insert	$O(1)$	$O(n)$	$O(\log n)$	$O(\log n)$
min	$O(n)$	$O(1)$	$O(1)$	$O(\log n)$
removeMin	$O(n)$	$O(1)$	$O(\log n)$	$O(\log n)$
remove	$O(1)$	$O(1)$	$O(\log n)$	$O(\log n)$
replaceKey	$O(1)$	$O(n)$	$O(\log n)$	$O(\log n)$
replaceValue	$O(1)$	$O(1)$	$O(1)$	$O(1)$

## Sorting: running times in practice (ms)

<i>Array Size</i>	<i>Selection Sort</i>	<i>Insertion Sort</i>	<i>Heap Sort</i>
1000	5,64	3,90	1,33
10000	80,45	25,97	3,85
50000	1661,01	348,15	12,44
100000	6792,87	1379,62	19,62
500000	164748,01	34541,45	87,24
1000000	670833,80	138192,96	248,09

~11 mins

~0.25 sec