

### Warm-up

**Problem 1.** Consider a hash table of size  $N > 1$ , and the hash function such that  $h(k) = k \bmod 2$  for every  $k$ . We insert a dataset  $S$  of size  $n < N$ . After that, what is the typical running times of GET for chaining and open addressing (as a function of  $n$ )?

**Solution 1.**  $\Theta(n)$  for both. Since the hash function maps all elements to only two values, 0 and 1, there are at least  $n/2$  elements hashed to the same value – for instance, 0. In the case of chaining, that means that at least  $n/2$  elements of  $S$  are now part of a linked list, and so doing a GET on any of those boils down to doing a linear search into a linked list of size  $\Theta(n)$ :  $\Theta(n)$ . In the open addressing case, similarly we have a sequence of at least  $n - 1$  contiguous positions corresponding to those elements: indeed, everything is mapped to either 0 or 1, so after the first 0 everything collides at 1 and every single following index until an empty slot is found (and if there is no element mapped to 0, then all  $n$  insertions collide at 1), and so we end up doing a search in an unsorted array of size  $\Theta(n)$ .

Upshot: the choice of hash function matters! Don't come up with your own, it may be a bad idea...

**Problem 2.** Suppose you are given a hash function  $h$  mapping 10-digit integers to integers in  $\{1, 2, \dots, 10000\}$ . Show that there is some dataset  $S$  of size 1,000,000 such that all keys of  $S$  are hashed to the *same* value.

**Solution 2.** There are  $10^{10}$  different 10-digit integers. By the Pigeonhole Principle, there exists some  $i \in \{1, 2, \dots, 10000\}$  such that  $|h^{-1}(i)| \geq \frac{10^{10}}{10000} = 10^6$  (where  $h^{-1}(i) \subseteq S$  is the set of keys  $k$  such that  $h(k) = i$ ). Let  $S$  be this set  $h^{-1}(i)$ .

Upshot: for every hash function  $h$ , there exists some dataset  $S = S(h)$  (depends on the hash function!) for which  $h$  is arbitrarily bad. So all we can ask for is that hash functions be good for *most* (i.e., "typical" for our applications) datasets, not *all*.

**Problem 3.** Work out the details of implementing cycle detection in cuckoo hashing based on the number of iterations of the eviction sequence.

**Solution 3.** There are  $2N$  entries in total, so if the cuckoo eviction process runs more than  $4N$  times, we are guaranteed to have a cycle: there can be  $2N$  shifts to move all entries to their alternative position and then another  $2N$  to move everything back to their initial positions. By adding a counter that is incremented every time we evict an entry, this is easily checked in  $O(1)$  time after every eviction.

**Problem 4.** Work out the details of implementing cycle detection in cuckoo hashing based on keeping a flag for each entry.

**Solution 4.** We augment the hash table entries with a boolean attribute called flag. Assume that at the start of the put routine all entries are unflagged (i.e., all flags are

set to false). Suppose we are trying to put a new element  $x$  into the hash table and that the element could store in hash entries  $i$  or  $j$ . To test if there exists an eviction path starting at, say,  $i$ , we design a recursive auxiliary routine  $\text{TEST-CHAIN}(x, i)$ .

This auxiliary routine first checks if  $i$  is empty; if that is the case, then the test is successful. If the entry is not empty, it checks if the entry is flagged; if that's the case the test is unsuccessful. Finally, if the entry is not empty and not flagged, let  $y$  be the item currently stored at  $i$  and  $k$  be the alternative entry for  $y$ . We flag entry  $i$ , and call  $\text{TEST-CHAIN}(y, k)$ . Regardless of the outcome (successful or unsuccessful) before returning we unflag the entry previously flagged so that all entries are unflagged at the start of the next put call.

If both  $\text{TEST-CHAIN}(x, i)$  and  $\text{TEST-CHAIN}(x, j)$  are unsuccessful, then it is not possible to put  $x$  into the has table. The running time is proportional to the length of the eviction chain tests.

Note that the unsuccessful eviction chain test could be much longer than the successful one. We could modify the algorithm to interleave the search for a chain so that we do work proportional to the length of the shortest successful chain (if one exists) but this would complicate the algorithm a great deal and would require us to keep two flags. Not really worth the effort.

### Problem solving

**Problem 5.** Design a sorted hash table data structure that performs the usual operations of a hash table with the additional requirement that when we iterate over the items, we do so in the order in which they were inserted into the hash table. Iterating over the items should take  $O(n)$  time where  $n$  is the number of items stored in the hash table. Your data structure should only add  $O(1)$  time to the standard put, get, and delete operations.

**Solution 5.** In addition to the hash table, we keep a doubly linked list of the entries and augment each entry in the hash table to also have a pointer to its position in the list.

When we need to put a new item, after inserting it into the hash table in the usual way, we add it to the end of the list and set the pointer of the entry in the hash table accordingly in  $O(1)$  time.

When we remove an item, after removing it from the hash table in the usual way, we use the pointer to the doubly linked list to remove it from there as well in  $O(1)$  time.

When we update an existing item, after updating it in the usual way, we move its position in the list to the end in  $O(1)$  time or we leave it in place depending on how we want to interpret this case; i.e., we care about the order of when the key of the item arrived or when the value of the item arrived.

Whenever we need to iterate over the entries, we use the linked list. Since iterating through all elements of a doubly linked list takes  $O(n)$  time, this satisfies our requirements.

**Problem 6.** Given an array with  $n$  integers, design an algorithm for finding a value that is most frequent in the array. Your algorithm should run in  $O(n)$  expected time.

**Solution 6.** Keep a hash table with  $2n$  entries using linear probing where the keys are the integers in the input array and the value is the frequency of the key.

We scan the integers in the array, updating their frequency in our hash table. That is, when processing  $k$ , we first try to get the entry for  $k$ . If there is no entry, we put  $(k, 1)$ ; otherwise, if there is already an entry  $(k, f)$  we updated with  $(k, f + 1)$ . At the end we do a scan of the table to find the integer with maximum frequency.

Given that the load factor of the hash table is  $\leq 1/2$ , the hash table operations take  $O(1)$  expected time. Finally, scanning the whole hash table to find the maximum frequency integer take  $O(n)$  time.

It is worth noting that we can avoid the scan of the hash table if we also keep track of the maximum frequency item we have seen so far. Every time we update the frequency of a number, we check if its frequency is larger than the maximum frequency so far, and if so, we update it.

**Problem 7.** A multimap is a data structure that allows for multiple values to be associated with the same key. The method `GET( $k$ )` should return all the values associated with key  $k$ . Describe an implementation where this method runs in  $O(1 + s)$  expected time, where  $s$  is the number of values associated with  $k$ .

**Solution 7.** (*Sketch*) Each entry, instead of having a single value, has a linked list of values associated with the key. When putting a new item  $(k, v)$  we add  $v$  to the list in the entry associated with  $k$ . When getting a key  $k$ , we find it in  $O(1)$  expected time and we go over the list in the entry associated with  $k$  (which takes  $O(s)$  time).

**Problem 8.** Suppose that you have a group of  $n$  people and you would like to know if there are two people that share a birthday. Design an  $O(1)$  time algorithm that given the information about the  $n$  people's birthdays, finds a pair that shares a birthday, or reports that no such pair exists.

**Solution 8.** Map the birthday to a number in  $[1, 365]$  and use a hash function on integers to build a hash table of size  $2 * 365$  using linear probing. We iterate over the elements in the list, for each person  $p$  we treat their birthday as a key  $k$  and try to get the entry with  $k$ . If there is no such key we add  $(k, p)$  to the hash table; otherwise, if we find the item  $(k, p')$  we have our pair  $p, p'$  of people sharing a birthday and we can stop.

If we scan the whole set of  $n$  people without finding a match, we report that none share a birthday.

For the time complexity, we note that if  $n > 365$  we are guaranteed to find a matching pair of people in the first 366 entries of the array so the algorithm terminates after  $O(1)$  iterations and each iteration takes  $O(1)$  time since the hash table has  $O(1)$  size.

**Problem 9.** In computational linguistics, texts (such as a book or an article) are modelled as a sequence of words. A  $k$ -gram is a sequence of  $k$  consecutive words.

A common task in language modelling requires that we compute the frequency of all  $k$ -grams that appear in the text.

Given a text with  $n$  words, design an  $O(n)$  expected time algorithm that computes the frequency of all  $k$ -grams that appear at least once in the text.

**Solution 9.** We use an approach similar Problem 6, but instead we use the  $k$ -grams in the text as our keys. To design a hash function for these keys we can use polynomial evaluation (since the order of the words of the  $k$ -gram matters).

Note that there can be at most  $n - k + 1$  distinct  $k$ -grams in the text, so keeping a hash table of size  $2n$  using linear probing yields the desired running time.

### Advanced problem solving

**Problem 10.** A hash table, and hashing in general, allow you to store and retrieve data quickly (in expectation, for typical data); the data structures never “make mistakes” (the result is always correct), but they do use some space: if each key takes  $b$  bits to store, a hash table will take  $O(nb)$  space to store  $n$  elements. In this problem, we will see a related data structure, the *Bloom filter*, which uses much less space while still providing efficient access and insertion: only  $O(n)$  space to store  $n$  elements (regardless of how many bits an element takes to encode)! This comes at a price: sometimes, the result of a query to the data structure will be wrong. However, when designed well, the frequency with which those mistakes occur is relatively low, and can be controlled.

Let’s start with what a Bloom filter is. For simplicity, we will only allow insertions (INSERT) and lookups (GET), but no deletions (they could be implemented, but this adds quite a bit of complexity to the data structure). The Bloom filter is an array  $A$  of size  $N$  (containing  $N$  bits, initialised to 0), along with  $m$  distinct hash functions  $h_1, \dots, h_m$ , each mapping the data universe  $U$  to  $\{1, 2, \dots, N\}$ . Here’s how it works:

- INSERT( $k$ ) evaluates the  $m$  hash functions on  $k$  and sets the bit of all  $m$  corresponding cells to 1.

```

1: function INSERT( $k$ )
2:   for all  $1 \leq i \leq m$  do
3:      $A[h_i(k)] \leftarrow 1$ 

```

- GET( $k$ ) evaluates the  $m$  hash functions on  $k$  and checks that all the  $m$  bits in the corresponding cells are equal to 1.

```

1: function GET( $k$ )
2:   for all  $1 \leq i \leq m$  do
3:     if  $A[h_i(k)] = 0$  then
4:       return False
5:   return True

```

That's all! In what follows, we will try to see what this does, how to analyse the performance, and see how to choose the parameter  $m$  (number of hash functions).

- a) What type of "mistakes" can GET make? Can it ever return "false" when the element is in the Bloom filter (*false negative*)? Can it ever return "true" when the element is not (*false positive*)?
- b) Assuming each hash function can be stored in  $O(1)$  space and takes  $O(1)$  time to evaluate, what is the space complexity of the Bloom filter? What are the time complexities of INSERT and GET?
- c) Let's analyse the error rate, that is, how frequently we would expect GET to make a mistake, "on average." In what follows, assume we inserted a dataset  $S$  of  $n$  elements into the Bloom filter. We will make the following (false, but convenient) assumption that we have perfectly random hash functions: the  $(h_i(k_j))_{i,j}$  are independent across keys  $k_j \in U$  and hash functions  $1 \leq i \leq m$ , and  $h_i(k)$  is uniformly distributed in  $\{1, 2, \dots, N\}$  for every  $i$  and every  $k$ :

$$\forall i, j, k, \Pr[h_i(k) = j] = \frac{1}{N} \quad (\text{Uniform hashing})$$

- 1) Fix any  $1 \leq i \leq N$ . After inserting  $n$  elements into our Bloom filter, what is the probability  $p_i$  that the  $i$ -th bit of our array  $A$  is set to 1?  
Let  $B := \frac{N}{n}$  be the average number of bits used per element. Using the approximation  $1 + x \approx e^x$  (very accurate for small  $x$ ), show that  $p_i \approx 1 - e^{-m/B}$ .
- 2) *Error rate*: What is the probability that, when calling GET( $k$ ) on a key which was *not* inserted (not part of the  $n$  keys from  $S$ ), the value returned is True?
- 3) Say you have a target per-key storage value  $B$  in mind:  $B = 8$  bits. What is the number of hash functions  $m$  you should use to minimise the probability of error?
- 4) For the setting  $B = 8$ , and the choice of  $m$  above, what is the error rate you should expect?
- 5) Let's use  $m = 6$  hash functions and explore the trade-off between space (parameter  $B$ ) and error rate – we could decide to use more space than 8 bits per element. What is the expected error rate if you increase  $B$  to 12 bits? 16? 32?

**Solution 10.**

- a) Only false positives, never any false negative. This is because once its corresponding bits are set to 1, the element will always be reported as present. But a non-present element might have its  $m$  bits set to 1 by several other elements. Here's an example with  $m = 3$ ,  $N = 10$ , and  $U = \{1, 2, 3, 4\}$ : consider the hash functions

	$h_1$	$h_2$	$h_3$
1	1	5	10
2	9	4	6
3	1	6	3
4	7	3	8

Say we insert  $S = \{1, 2, 4\}$ : the corresponding bits set to 1 will be (indexing starting at 1)

$$1 \rightsquigarrow A[1], A[5], A[10]$$

$$2 \rightsquigarrow A[9], A[4], A[6]$$

$$4 \rightsquigarrow A[7], A[3], A[8]$$

Now, when calling `GET(3)`, we will check if the bits  $A[1], A[6], A[3]$  are all equal to 1. And they all are, so `GET(3)` will return `True` even though 3 was *not* inserted.

- b) There are  $m$  hash functions, each taking constant space; and the hash table itself takes  $O(N)$  space, hence  $O(m + N)$  space suffices. The time complexities are all  $O(m)$ , since they loop over (at most)  $m$  hash functions and, for each, perform constant time operations (computing the value of the current hash function, and accessing a position of the hash table).

c)

- 1) Since we made the assumption of truly uniform hashing, the probability that, for any fixed element  $k$  inserted, the  $i$ -th bit is *not* set to 1 by the  $j$ -th hash function is equal to  $1 - 1/N$ . By independence, since we have  $m$  hash functions and  $n$  elements, the probability that the  $i$ -th bit is *not* set to 1 is equal to  $(1 - 1/N)^{mn}$ , and so

$$p_i = 1 - \left(1 - \frac{1}{N}\right)^{mn} \approx 1 - e^{-\frac{nm}{N}} = 1 - e^{-\frac{m}{B}}$$

- 2) For this to happen, we need *all*  $m$  bits  $h_1(k), \dots, h_m(k)$  to be set to 1. By the previous question and our independence assumption, this happens with probability

$$p_1 \times \dots \times p_m = \left(1 - e^{-\frac{m}{B}}\right)^m$$

- 3) Either eyeball it by a plot, or use calculus (differentiate  $\left(1 - e^{-\frac{m}{8}}\right)^m$  with respect to  $m$ ). You might want to use <https://www.wolframalpha.com/>. In detail: letting  $f(x) = (1 - e^{-x/8})^x$ , we want to minimise  $f$ . Differentiating, you can check that

$$f'(x) = f(x) \left( \frac{x}{8} \cdot \frac{1}{e^{x/8} - 1} + \ln(1 - e^{-x/8}) \right)$$

and, since  $f(x) > 0$  for all  $x > 0$ ,  $f'(x) = 0$  if, and only if,

$$\frac{x}{8} \cdot \frac{1}{e^{x/8} - 1} + \ln(1 - e^{-x/8}) = 0.$$

Going further to argue that there is exactly one solution requires more calculus and is not very interesting, but you can check that plugging  $x = 8 \ln 2$  in the left-hand side does evaluate to 0:  $f$  is minimised for  $x = 8 \ln 2 \approx 5.6$ .

The right answer is therefore  $m = 6$  (the function is minimised for  $m \approx 5.6$ , and we need an integer). In general, one can derive the answer (again, based on the above approximations and assumptions, which are actually quite well supported in practice) to be  $m = \lceil (\ln 2)B \rceil$ . See, e.g., the above computation replacing 8 by  $B$ , or this computation on WolframAlpha.

- 4) We have  $(1 - e^{-6/8})^6 \approx 0.0216$ , so the expected false positive rate when calling GET is roughly 2.16%.
- 5) The corresponding values are 0.37%, 0.09%, and... 0.0025%. The error rates decrease *exponentially* in  $B$  (for fixed  $m, n$ ): see this plot.