

Database-backed applications

ISYS2120 Data and Information Management

Prof Alan Fekete

University of Sydney

Acknowledge: slides from Uwe Roehm and Alan Fekete, and from the materials associated with reference books (c) McGraw-Hill, Pearson

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**). The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.

Database-backed software

- Most of the software used for life and business needs to use data that is found in a database
 - Often, user activities will also modify the data in the database
- A few examples
 - E-commerce (inventory, purchase, order status, order history)
 - Entertainment (catalogue, preferences, history)
 - Social media (posts, community connections)
 - Transport (routes, timetables, current status)
- The software that provides these functionalities, needs to access one or more dbms

Software and queries

- Previously in isys2120, we covered how to write SQL to extract and modify data
- A human typed the query into a query window and ran it against the database, and observed the output
- End-users can't be expected to know SQL, nor usually to have accounts and access on all the dbms needed for modern life
- Instead, the user invokes some software, and the software submits queries to dbms, gets result table, and displays information to the user
- This lecture is about the structure of that software

Data-intensive Systems

- Three types of functionality (often placed in separate layers of code):

Presentation Logic

- Input – keyboard/mouse/gestures
- Output – monitor/printer/screen

GUI Interface

Processing Logic

- Business rules
- I/O processing

Procedures, functions, programs

Data Management

(Storage Logic)

- data storage and retrieval


DBMS activities

- The system architecture determines whether these three components reside on a single computing system (1-tier) or whether they are distributed across several tiers

Presentation layer

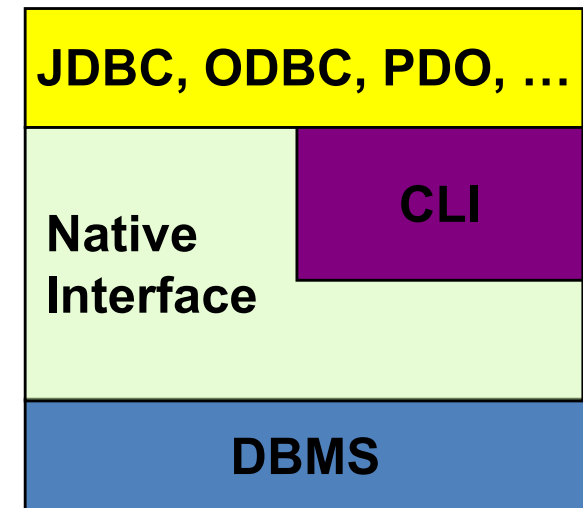
- Often, web browser is used as GUI for end-user
 - Available on all devices!
- Application needs to have code that lays out the web pages and provides navigation between them
- Mobile devices may provide an app that directly works with the gestures etc of the device
 - And is targeted for the small screen size
- Lab08 and Asst3 will work with web interface
 - Flask library of Python, to construct a web-server that also runs business logic and data management

SQL in Application Code

- SQL commands can be called from within a *host language* (such as Python or Java) program.
 - Must include a statement to *connect* to the right database.
 - SQL statements can refer to *host variables* (including special variables used to return status).
- Two main integration approaches:
 - **Statement-level interface (SLI)**
 - Embed SQL in the host language (Embedded SQL in C, SQLJ)
 - Application program is a mixture of host language statements and SQL statements and directives
 - A special compiler must deal with both aspects
 - **Call-level interface (CLI)**  This is what Flask uses
 - Create special API to call SQL commands (JDBC, ODBC, Python, ...)
 - SQL statements are passed as arguments to host language (library) procedures / APIs
 - Standard programming language compiler, and program is combined with a library that supports the API

Call-level Interfaces and Database APIs

- Program can invoke methods/procedures in a library with database calls (API)
 - Pass SQL strings from language, present result sets in language-friendly way
 - Supposedly DBMS-neutral
 - a “driver” executes the calls and translates them into DBMS-specific code
 - database can be across a network
- Several Variants
 - **SQL/CLI**: “SQL Call-Level-Interface”
 - Part of the SQL-92 standard;
 - “The assembler under the APIs”
 - **ODBC**: “Open DataBase Connectivity”
 - Side-branch of early version of SQL/CLI
 - Enhanced to: **OLE/db**, and further **ADO.NET**
 - **JDBC**: “Java DataBase Connectivity”
 - Java standard
 - **PDO**
 - Persistency standard for PHP Data Objects



Whose privileges when code is run?

- Many databases are accessed indirectly
 - End-user does not write and submit SQL, but rather runs a program that (team of) coders have written to perform useful activity
 - Eg a student changes their address through MyUni
- The program can do lots of checking of whether access is appropriate, before sending SQL to dbms
 - Also the program can filter or summarise data, so user does not see everything the program gets from the dbms
- The program may run with its own appropriate level of privilege (or that of the coders), rather than from the end-user who is the source of request
 - Indeed, the end-user may not have a dbms account at all
- Often, the program has quite a lot of privilege, but this is risky if there are mistakes in the code, or if an attacker can obtain the program's credentials [eg if program uses a password which is stored somewhere, and leaked]

PYTHON DB-API2

a Call-Level API Example

Python

- Python features extensive standard library (modules)
 - Special functionality supported by variety of optional 3rd-party modules
 - For database connectivity, several database-specific python modules
 - e.g. psycopg or pg8000 (PostgreSQL) or cx_oracle (Oracle)
 - <https://pypi.org/project/pg8000/>
 - For dynamic websites:
 - several framework available; in lab8, asst3 we will use Flask
 - Allows to define template pages of html with embedded python code

Python Database API Specification (DB-API)

- DB-API 2.0 was released April 1999
- Defines common functions and API for access modules to different database systems
 - Module API; Connection and Cursor interface definitions
- Works as a generic as a **database abstraction layer**
 - Generic driver model to connect to different database engines via the same API

- URLs to learn more:

<https://pypi.org/project/pg8000/>

<https://www.python.org/dev/peps/pep-0249/>

<http://initd.org/psycopg/docs>

http://www.tutorialspoint.com/postgresql/postgresql_python.htm

<https://wiki.python.org/moin/DatabaseProgramming>

<https://wiki.python.org/moin/UsingDbApiWithPostgres>

Python DB-API Example

```
import pg8000

try:
    # connect to the database
    conn = pg8000.connect(database="postgres",user="test",password="secret")

    # prepare to query the database
    curs = conn.cursor()

    # execute a parameterised query
    unit_of_study = "ISYS2120"
    curs.execute("""SELECT name
                    FROM Student NATURAL JOIN Enrolled
                    WHERE uos_code = %(uos)s""", {'uos': unit_of_study} )

    # loop through the resultset
    for result in curs:
        print (" student: " + result[0])

    # clean up
    curs.close()
    conn.close()

except Exception as e: # error handling
    print("SQL error: unable to connect to database or execute query")
    print(e)
```

Core tasks with SQL Interfaces

(1) Establishing a database connection

(2) Static vs. Dynamic SQL

(3) Parameterized SQL and mapping of domain types to data types of host

- ▶ Concept of *host variable*

- ▶ How to treat *NULL* values?

(4) Impedance Mismatch:

- ▶ SQL operates on sets of tuples

- ▶ Host languages like C do not support a set-of-records abstraction, but only a one-value-at-a-time semantic

- ▶ Solution: *Cursor Concept*

Iteration mechanism (loop) for processing a set of tuples

(5) Error handling

(1) DB Connections from Python

- Session with PostgreSQL started by creating a connection
- Two Variants:
 - Connect with keyword arguments

```
conn =  
pg8000.connect(host='...', database='...', user='X', password='...')
```
 - Connect with a Data Source Name (*DSN*) string of the form
"host=X dbname=Y user=U password=P"
For example for PostgreSQL:

```
conn = pg8000.connect(  
"host=postgres.usyd.edu.au dbname=unidb user=U  
password=secret")
```

Python Database Connection Modules

- Python support for variety of DBMSs
 - MySQL (module: MySQLdb)
 - PostgreSQL (module: pg8000 or psycopg2)
 - Oracle (module: cx_oracle)
 - IBM DB2 (module: ibm_db)
 - SQL Server (module: pymssql)
 - sqlite (module: sqlite3)
 - ...
 - DSN syntax and additional DB parameters vary for each driver
 - Check manuals...

Note:

db modules need to be installed first as part of the Python installation...

- Example for Oracle:

```
dsnStr = cx_oracle.makedsn("oracle10g.it.usyd.edu.au",1521,"ORCL")  
conn = cx_oracle.connect(user="myuser",password="mypass",dsn=dsnStr)
```


pg8000 Connection Simple Example

```
import pg8000

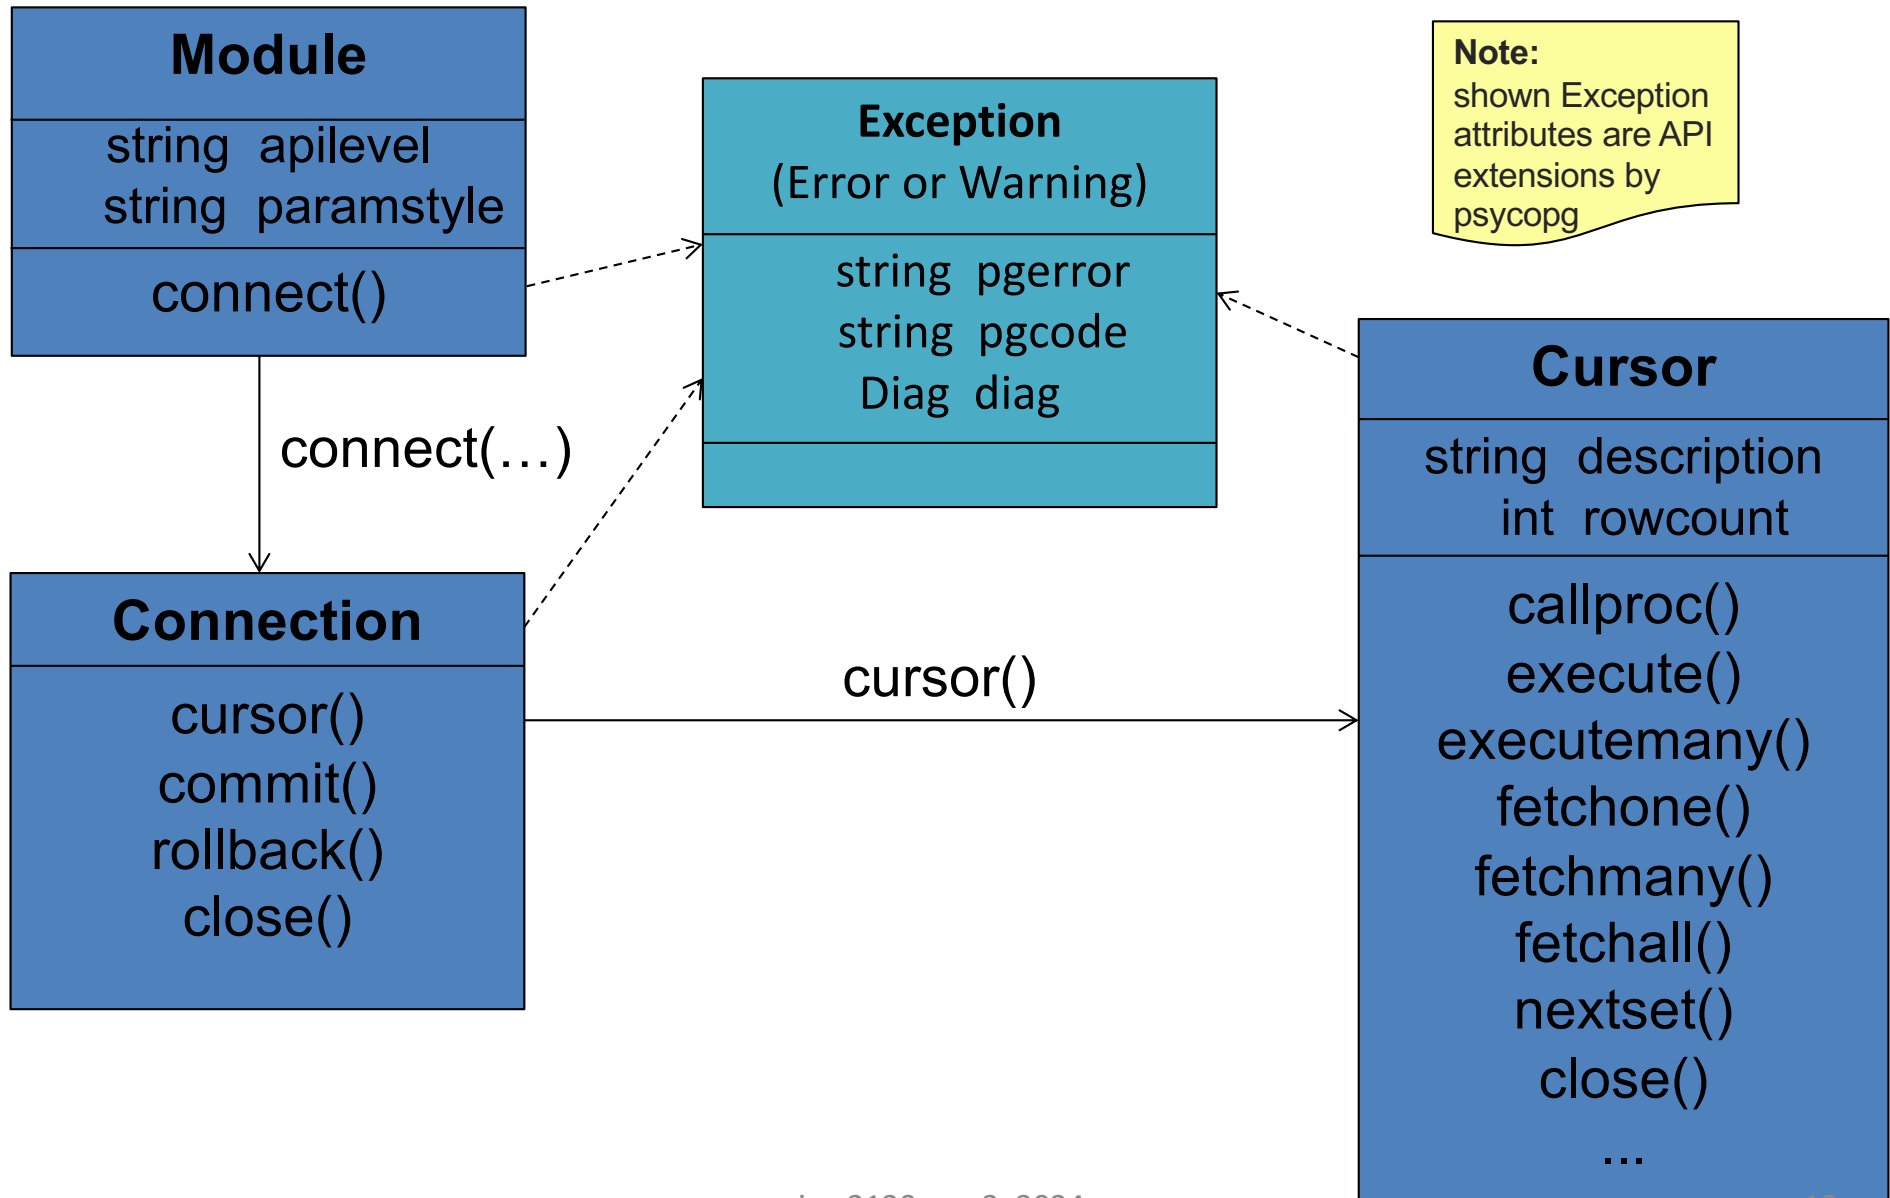
# connect to the database
try:
    conn = pg8000.connect(database='foo', user='dbuser', password='pwd')
except:
    print("unable to connect to database")

# query database
curs = conn.cursor()
try:
    curs.execute("SELECT name FROM Student WHERE studID=4711")
except:
    print("unable to execute query")

... Do Actual Work ...

# cleanup
curs.close()
conn.close()
```

Python DB-API 2.0 Objects



Python Database API Interfaces

- Connection Management
 - `pg8000.connect()` connects to a database
 - `conn.cursor()` creates a cursor object for query execution
- Start SQL statements
 - `cursor.execute()` for static SQL, and also parameterized SQL queries
 - `cursor.callproc()` for executing a stored procedure including parameters
- Result retrieval
 - `cursor.fetchone()` retrieves next row of a result or **None** when no more data
 - `cursor.fetchall()` retrieves the whole (remaining) result set, and returns it as a list of tuples
- Transaction control
 - `conn.commit()` successfully finishes (commits) current transaction
 - `conn.rollback()` aborts current transaction
- Error Handling
 - Via standard exception handling of Python

Side Note on DB Connections

- Establishing a database connection takes some time...
 - Network communication, memory allocation, dbs authorization
- So do this only once in your program
 - ... **not over and over** for individual SQL queries
- Modern, multi-threaded applications will typically want to have a pool of connections that are re-used
 - Might be handled by your runtime library (that's what happens in Python)
 - But for, e.g., Java programs better be mindful of connection costs!

(2) Static vs. Dynamic SQL

- SQL constructs in an application can take two forms:
 - Static SQL statements:
Useful when SQL query is fully known at compile time
 - no parameters are allowed in the query string
 - only useful in context of compiled languages such as C
 - Dynamic SQL statements:
Application determines SQL statements at *run time* as values of host language variables that are manipulated by directives.
 - Challenge: Python is not a compiled language;
everything in Python/pg8000 is by definition dynamic SQL...
 - This means we have to be careful on how we construct any query
and in particular how parameters are passed to the database

DB-API: Executing SQL Statements

- Three different ways of executing SQL statements:
 - ▶ ***cursor.execute(sql)*** semi-static SQL statements
 - ▶ ***cursor.execute(sql,params)*** parameterized SQL statements
 - ▶ ***cursor.callproc(call,args)*** invoke a stored procedure in DBMS
 - ▶ ***cursor.executemany(sql,seq_of_params)*** repeatedly executes parameterized SQL statements

- In DB-API 2.0,
 - Need to create new cursor and re-issue SQL statement each time when parameters change – or if possible use ***executemany()***
 - Some other APIs offer “prepared statements” – parsed and optimized once in the dbms, then re-executed over and over with different parameters

Python DB-API with fixed SQL

- Simplest way to execute an unchanging SQL query:

```
import pg8000

try:
    # connect to the database
    conn = pg8000.connect(database='foo', user='dbuser', password='pwd')

    # query database
    curs = conn.cursor()
    curs.execute("SELECT name FROM Student WHERE studID=4711")
    result = curs.fetchone()
    print(result)

    # cleanup
    curs.close()
    conn.close()

except:
    print("unable to connect to db or to execute query")
```

DB-API: Batch Insert Example

- Example: executing batch INSERT statements

```
import pg8000

try:
    # connect to the database
    conn = pg8000.connect(database='foo', user='dbuser', password='pwd')

    # prepare list of insert values (3 students enrolling in ISYS2120)
    params = [(4711, 'ISYS2120'), (4712, 'ISYS2120'), (4713, 'ISYS2120')]

    # execute INSERT statement batch
    curs = conn.cursor()
    curs.executemany("INSERT INTO Enrolled VALUES (%s,%s)", params)
    conn.commit() # cf. next week on transactions

    # cleanup
    curs.close()
    conn.close()

except:
    print("unable to connect to db or to execute query")
```


DB-API: Parameterized Queries

- Two (safe) approaches for passing query parameters:
(because `execute()` will do any necessary escaping / conversions for parameter markers)

1. Anonymous Parameters

```
studid = 12345
cursor.execute(
    "SELECT name FROM Student WHERE sid=%s",
    (studid,) )
```

This comma is no mistake, but needed with single parameters

parameter marker

2. Named Parameters

```
studid = 12345
cursor.execute(
    "SELECT name FROM Student WHERE sid=%(sid)s",
    {'sid': studid} )
```

named parameter marker

(3) Parameterized SQL & Host Variables

- Data transfer between DBMS and application
- Mapping of SQL domain types to data types of host language
- Python DB-API:
 - Host variables are normal *dynamically typed* Python variables; automatic conversion to/from SQL types done by pg8000 in

`execute()`:

```
studid = 12345
stmt = cursor.execute(
    "SELECT name FROM Student WHERE sid=%s",
    (studid,) )
```

- Note: in statement-level interface such as ESQL/C:
Host variables must be declared before usage

```
EXEC SQL BEGIN DECLARE SECTION;
    int  studid = 12345;
    char sname[21];
EXEC SQL END DECLARE SECTION;
```

*Variables
shared by host
and SQL*

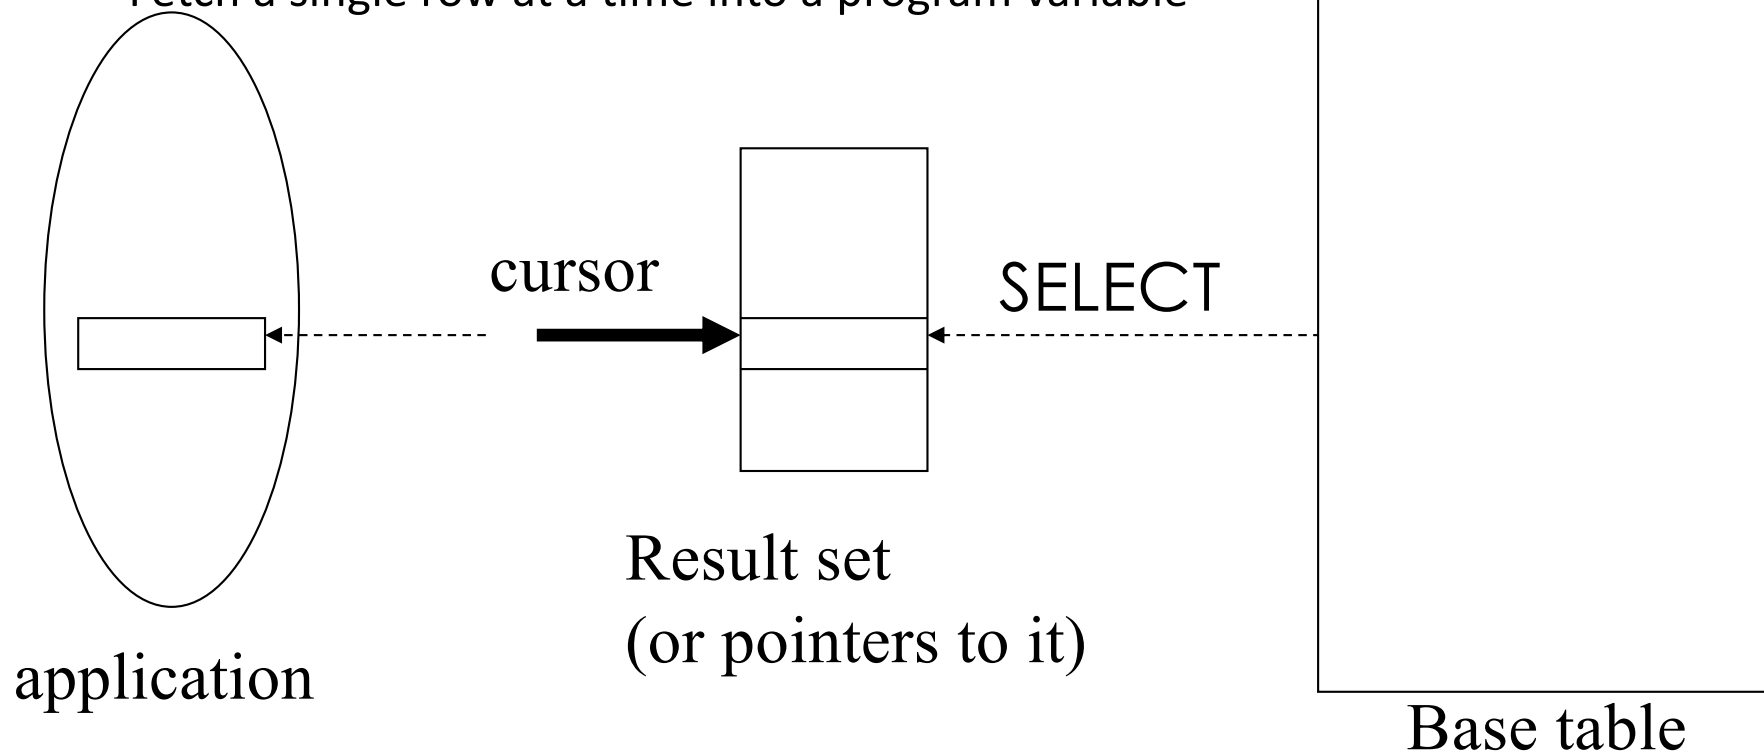
Specifying Date/Time: Type Objects

- Providing date or time values is very database specific with different database configurations requiring particular formats
 - eg. "26.04.2016" (AU) versus "04/26/2014" (US)
- DB-API has helper objects to safely convert database types
 - **Date** (*year , month , day*)
 - **Time** (*hour , minute , second*)
 - **Timestamp** (*year , month , day , hour , minute , second*)
 - **DateFromTicks** (*ticks*)
 - **TimeFromTicks** (*ticks*)
 - **TimestampFromTicks** (*ticks*)
 - **Binary** (*string*)

(4) Buffer Mismatch Problem

(also called: Impedance Mismatch)

- **Problem:** SQL deals with tables (of arbitrary size); host language program often deals with fixed size data types (this is less true with Python which has collections)
 - How is the application to allocate storage for the result of a SELECT statement?
- **Solution:** Cursor concept
 - Fetch a single row at a time into a program variable



Mapping of Sets: Cursor Concept

- ***Result set*** – set of rows produced by a SELECT statement
- ***Cursor*** – pointer to a row in the result set.
- Cursor operations:
 - *Declaration*
 - *Open* – execute SELECT to determine result set and initialize pointer
 - *Fetch* – advance pointer and retrieve next row (Python: fetchone() call)
 - *Close* – deallocate cursor

Cursor in Python– via Cursor Interface

■ Cursor concept with Python/psycopg:

```
cursor = conn.cursor()
cursor.execute("SELECT title, name, address FROM Emp")
row = cursor.fetchone()
while row is not None:
    print(row)
    row = cursor.fetchone()
cursor.close()
```

■ Cursor objects are iterable, so shorter form is:

```
cursor = conn.cursor()
cursor.execute("SELECT title, name, address FROM Emp")
for row in cursor:
    data = row[0] + "\t" + row[1] + "\t" + row[2] + "\n"
    print(data)
cursor.close()
```

You address result
columns by position

isys2120 sem2, 2024

Cursor in Python – fetchAll()

- Fetchall() returns a Python list with all the result rows

- Good for small results

```
curs.execute("SELECT title, name, address FROM Emp")
resultset = curs.fetchall()
curs.close()
for row in resultset:
    print(row)
```

- ▶ just be mindful that this will be **memory intensive** for large results

Dictionary Cursors

- By default, pg8000 returns tuples with `fetch()` / `fetchall()`
 - fields can only be addressed positionally
- As an extension, pg8000 also supports dictionary cursors
 - Result is now a dictionary (associative array) which each field being named by the attribute names from the database schema

```
import pg8000
from pg8000.extras import RealDictCursor
...
curs = conn.cursor(cursor_factory=RealDictCursor)
curs.execute("SELECT title, name, address FROM Emp")
for row in curs:
    data = row['title'] + "\t" + row['name'] + "\n"
    print(data)
curs.close()
```


NULL Handling in Python

- Remember: In SQL there is a special indication NULL used for unknown or inapplicable value of a column
 - Null value is not the same as 0 nor empty string
- In Python this shows as **None**:

```
cursor.execute("SELECT gender FROM Student ...")
result = cursor.fetchone()
if result[0] is None:
    # null value
else:
    # no null value
```

- Other languages require a special *indicator variable*. Eg. C:

```
EXEC SQL select gender into :gender:indicator
from Student where

sid=4711;
if ( indicator == -1 )
{ /* null value */ }
else
{ /* no null value */ }
```

Testing for Variable Exists / is None

- In Python, to check for existence of a variable versus whether it has None as value:

Ensure variable is defined

```
try:
```

```
    x
```

```
except NameError:
```

```
    x = None
```

Test whether variable is defined to be None

```
if x is None :
```

```
    some_fallback_operation()
```

```
else:
```

```
    some_operation(x)
```

[Source: <http://code.activestate.com/recipes/59892/>]

(5) Error Handling

- Multitude of potential problems
 - No database connection or connection timeout
 - Wrong login or missing privileges
 - SQL syntax errors
 - Empty results
 - NULL values
 - ...
- Hence always check database return values,
- Provide error handling code, resp. exception handlers
- Gracefully react to errors or empty results or NULL values
- **NEVER show database errors to end users**
 - Not only bad user experience, but huge security risk, because database errors can report table names, stack trace showing lines of code, etc

Error Handling with Python DB API

- Error handling via normal exception mechanism of Python
 - Errors and warnings are made available as Python exceptions
 - Warning raised for warnings such as data truncation on insert, etc
 - Error exception raised for various db-related errors
- psycopg has an API extension:
 - Exception attributes for detailed SQL error codes and messages
 - **pgerror** string of the error message returned by backend
 - **pgcode** string with the **SQLSTATE** error code returned by backend

Example:

```
try:
    conn = pg8000.connect(database="postgres", user="test", password="secret")

# error handling
except OperationalError as e:
    print("unable to connect to database")
except Exception as e:
    print("Error when querying database")
    print(e)
```

Note: please do not directly print SQL exceptions ;)

Exception Hierarchy of Python DB API

- The complete **Exception** inheritance hierarchy for the Python DB API is as follows:

```
StandardError
|__ Warning
|__ Error
|__ InterfaceError
|__ DatabaseError
|__      DataError
|__      OperationalError
|__      IntegrityError
|__      InternalError
|__      ProgrammingError
|__      NotImplementedError
```

[cf. <http://initd.org/psycopg/docs/module.html?highlight=connect#exceptions>]

Reprise: Example of Python DB-API

```
import pg8000
```

```
try:
```

```
# connect to the database
```

```
conn = pg8000.connect(database="postgres", user="test", password="secret")
```

```
# prepare to query the database
```

cursor concept

```
curs = conn.cursor()
```

```
# execute a parameterised query
```

```
unit_of_study = "ISYS2120"
```

```
curs.execute("""SELECT name  
                FROM Student NATURAL JOIN Enrolled  
                WHERE uos_code = %(uos)s""", {'uos': unit_of_study})
```

*dynamic SQL with
safe parameter passing*

```
# loop through the resultset
```

```
for result in curs:
```

```
    print (" student: " + result[0])
```

```
# clean up
```

```
curs.close()
```

```
conn.close()
```

```
# error handling
```

```
except OperationalError as e:
```

```
    print("unable to connect to database")
```

```
except Exception as e:
```

```
    print("Error when querying database")
```

```
    print(e)
```

error handling

Stored Procedures (Server-side application logic)

- Recall: these run application logic within the database server
 - Included as schema element (stored in DBMS)
 - Invoked by the application
- Advantages:
 - Central code-base for all applications
 - Improved maintainability
 - Additional abstraction layer (programmers do not need to know the schema)
 - Reduced data transfer
 - Less long-held locks
 - DBMS-centric security and consistent logging/auditing (important!)
- Note: although referred to as procedures, can also be functions

Python DB-API: Calling Stored Procedures

- **Cursor** objects have an explicit `callproc()` method
 - `cursor.callproc()` makes the OUT parameters available as resultset
- Example:

```
CREATE FUNCTION test(input VARCHAR,OUT output
VARCHAR) AS $$
BEGIN
    output := UPPER(input);
END $$ LANGUAGE plpgsql;
```

```
import pg8000
conn = pg8000.connect(...)
curs = conn.cursor()
input = "foo bar"
curs.callproc("test", [input] )
output = curs.fetchone()
print(output[0])
```

Pass all IN parameters
as a list in order of the
function declaration

OUT parameters are
returned as resultset

Language support of Stored Procedures

- Programming language virtual machine is often ‘integrated’ with DBMS
 - E.g. Java with Oracle
 - .Net CLR with IBM, Oracle, and SQL Server
 - PostgreSQL: Supports several scripting languages such as perl etc.
- But degree of integration differs heavily
 - If language VM is in a different process from dbms, then performance often suffers

Summary

- **Understand core issues for db-backed development**
 - Data and type conversion: *Host Variables*
 - NULL value semantic: *Indicator variables* and testing methods
 - Impedance Mismatch: *Cursor Concept*
 - *Dynamic versus static SQL: Passing parameters in and out*
- **Database APIs**
 - After lab08, you should be able to work with small Python db-accessing programs
- **Server-side database programming**
 - How to use stored procedures to run code inside a DBMS
 - e.g. with PostgreSQL's pl/pgsql
 - Modern database engines provide virtual machine environments to run external code near the data

References

- Silberschatz/Korth/Sudarshan(7ed)
 - Chapter 5.1, 9.1, 9.2, 9.4.2, 9.6, 9.7, 9.8
 - Ch 9 also covers other technologies, including Servlets, JSP, PHP, Django, Hibernate.

Also

- Kifer/Bernstein/Lewis(complete version, 2ed)
 - Chapter 8.1, 8.2, 8.4, 26.7
 - Ch 8 also covers other technologies, including JDBC (quite similar to DB-API), SQLJ
- Ramakrishnam/Gehrke(3ed)
 - Chapter 6.1, 7.5, 7.7, case study in 7.8
 - Chs 6,7 also cover other technologies, including JDBC, SQLJ
- Garcia-Molina/Ullman/Widom(complete book, 2ed)
 - Chapter 9.1-9.5
 - Ch 9 also covers JDBC and PHP

Database Documentation:

- Python DB-API: <http://initd.org/psycpg/docs/>
<https://wiki.python.org/moin/UsingDbApiWithPostgres>
http://www.tutorialspoint.com/postgresql/postgresql_python.htm